

# Fundamentals of Machine Learning (1618003)

Semester-1

Post Graduate Diploma in Data Science

Prepared By:

Krima Doshi(20137680015)

Course Coordinator:

Prof. Mahesh Panchal

Designation

GTU-GSET



Gujarat Technological University  
Graduate School of Engineering and Technology

December, 2020 to March, 2021

Assignment 2

March 24, 2021

# Contents

<b>1</b>	<b>Project Aim</b>	<b>3</b>
<b>2</b>	<b>Data description</b>	<b>4</b>
<b>3</b>	<b>Work Environment</b>	<b>5</b>
<b>4</b>	<b>Algorithm Selection</b>	<b>9</b>
<b>5</b>	<b>Neural Networks</b>	<b>10</b>
5.1	Preprocessing . . . . .	10
5.2	Model Preparation . . . . .	11
5.2.1	NN with keras . . . . .	11
5.2.2	DNN with keras . . . . .	12
5.3	Model Testing . . . . .	13
<b>6</b>	<b>Linear Regression</b>	<b>15</b>
6.1	Preprocessing . . . . .	15
6.2	Model Preparation . . . . .	15
6.3	Model Testing . . . . .	15
<b>7</b>	<b>Conclusion</b>	<b>17</b>

# List of Figures

1	Dataset . . . . .	9
2	Example of a Single Layer Neural Network . . . . .	10
3	Example of a Multi Layer Neural Network . . . . .	10
4	Normalized input . . . . .	11
5	MSE for NN 0.1 learning rate and epoch=10 . . . . .	13
6	MSE for DNN 0.1 learning rate and epoch=10 . . . . .	14
7	MSE for DNN 0.01 learning rate and epoch=100 . . . . .	14
8	Example of a linear regression model fitting . . . . .	15
9	Linear Regression params and error . . . . .	16

# 1 Project Aim

Prediction of the release year of a song from audio features. Songs are mostly western, commercial tracks ranging from 1922 to 2011, with a peak in the year 2000s.

## 2 Data description

Link to dataset: <https://archive.ics.uci.edu/ml/datasets/YearPredictionMSD>

Number of Instances: 515345

Number of Attributes: 90

Missing Values: None

Associated Tasks: Regression

Training Data: 463,715

Testing Data: 51,630

90 attributes

12 = timbre average

78 = timbre covariance

The first value is the year (target), ranging from 1922 to 2011. Features extracted from the 'timbre' features from The Echo Nest API. We take the average and covariance over all 'segments', each segment being described by a 12-dimensional timbre vector.

Timbre are the ways used to describe the sound, so words such as Light, Flat, Smooth, Smoky, Breathy, Rough, and so on are what you use to distinguish one sound from another. How you recognize the different sounds or voices you hear is attributed to the timbre.

Echo Nest classifies the timbre as: the first dimension represents the average loudness of the segment; second emphasizes brightness; third is more closely correlated to the flatness of a sound; fourth to sounds with a stronger attack; etc

### 3 Work Environment

Anaconda, jupyter-notebook, python 3.9, tensorflow, keras and scikit-learn are the major modules used.

Linear Regression has been used with a 8GB computer on a 64 bit OS with NVIDIA graphics card.

Neural Network has been applied on a 64 GB RAM GPU on a 64 bit OS with Quadro P400/PCIe/SSE2.

You can setup the environment on anaconda with the following environment.yml file:

```
1 name: mlenv
2 channels:
3   - pytorch
4   - conda-forge
5   - anaconda
6   - defaults
7 dependencies:
8   - _tflow_select=2.1.0=cpu
9   - absl-py=0.12.0=py38h06a4308_0
10  - aiohttp=3.6.3=py38h7b6447c_0
11  - anyio=2.2.0=py38h578d9bd_0
12  - argon2-cffi=20.1.0=py38h25fe258_2
13  - astunparse=1.6.3=py_0
14  - async-timeout=3.0.1=py38_0
15  - async-generator=1.10=py_0
16  - attrs=20.2.0=py_0
17  - babel=2.9.0=pyhd3deb0d_0
18  - backcall=0.2.0=pyh9f0ad1d_0
19  - backports=1.0=py_2
20  - backports.functools_lru_cache=1.6.1=py_0
21  - blas=1.0=mkl
22  - bleach=3.3.0=pyh44b312d_0
23  - blinker=1.4=py38_0
24  - brotlipy=0.7.0=py38h7b6447c_1000
25  - c-ares=1.17.1=h27cfd23_0
26  - ca-certificates=2020.12.5=ha878542_0
27  - cachetools=4.1.1=py_0
28  - certifi=2020.12.5=py38h578d9bd_1
29  - cffi=1.14.3=py38he30daa8_0
30  - chardet=3.0.4=py38_1003
31  - click=7.1.2=py_0
32  - coverage=5.3=py38h7b6447c_0
33  - cryptography=3.1.1=py38h1ba5d50_0
34  - cudatoolkit=10.1.243=h6bb024c_0
35  - cudnn=7.6.5=cuda10.1_0
36  - cupi=10.1.168=0
37  - cyclical=0.10.0=py_2
38  - cython=0.29.21=py38he6710b0_0
39  - dbus=1.13.18=hb2f20db_0
40  - decorator=4.4.2=py_0
41  - defusedxml=0.7.1=pyhd8ed1ab_0
```

```

42 - entrypoints=0.3=pyhd8ed1ab_1003
43 - expat=2.2.10=he6710b0_2
44 - faiss-gpu=1.7.0=py3.8_hf4b0e32_0_cuda10.1
45 - fontconfig=2.13.1=he4413a7_1000
46 - freetype=2.10.4=h7ca028e_0
47 - gast=0.4.0=py_0
48 - glib=2.67.4=h36276a3_1
49 - google-auth=1.22.1=py_0
50 - google-auth-oauthlib=0.4.1=py_2
51 - google-pasta=0.2.0=py_0
52 - grpcio=1.36.1=py38h2157cd5_1
53 - gst-plugins-base=1.14.0=hbbd80ab_1
54 - gstreamer=1.14.0=h28cd5cc_2
55 - h5py=2.10.0=py38hd6299e0_1
56 - hdf5=1.10.6=hb1b8bf9_0
57 - icu=58.2=hf484d3e_1000
58 - idna=2.10=py_0
59 - importlib-metadata=2.0.0=py_1
60 - intel-openmp=2020.2=254
61 - ipykernel=5.5.0=py38h81c977d_1
62 - ipython=7.21.0=py38h81c977d_0
63 - ipython_genutils=0.2.0=py_1
64 - jedi=0.18.0=py38h578d9bd_2
65 - jinja2=2.11.3=pyh44b312d_0
66 - joblib=0.17.0=py_0
67 - jpeg=9d=h36c2ea0_0
68 - json5=0.9.5=pyh9f0ad1d_0
69 - jsonschema=3.2.0=pyhd8ed1ab_3
70 - jupyter-packaging=0.7.12=pyhd8ed1ab_0
71 - jupyter_client=6.1.12=pyhd8ed1ab_0
72 - jupyter_core=4.7.1=py38h578d9bd_0
73 - jupyter_server=1.4.1=py38h578d9bd_0
74 - jupyterlab=3.0.12=pyhd8ed1ab_0
75 - jupyterlab_pygments=0.1.2=pyh9f0ad1d_0
76 - jupyterlab_server=2.3.0=pyhd8ed1ab_0
77 - keras=2.4.3=py_0
78 - keras-preprocessing=1.1.2=pyhd3eb1b0_0
79 - kiwisolver=1.3.1=py38h82cb98a_0
80 - ld_impl_linux-64=2.33.1=h53a641e_7
81 - libedit=3.1.20191231=h14c3975_1
82 - libfaiss=1.7.0=hbff11de5_0_cuda10.1
83 - libffi=3.3=he6710b0_2
84 - libgcc-ng=9.1.0=hdf63c60_0
85 - libgfortran-ng=7.3.0=hdf63c60_0
86 - libpng=1.6.37=h21135ba_2
87 - libprotobuf=3.14.0=h8c45485_0
88 - libsodium=1.0.18=h36c2ea0_1
89 - libstdcxx-ng=9.1.0=hdf63c60_0
90 - libtiff=4.0.10=hc3755c2_1005
91 - libuuid=2.32.1=h14c3975_1000
92 - libxcb=1.13=h14c3975_1002
93 - libxml2=2.9.10=hb55368b_3
94 - lz4-c=1.9.2=he1b5a44_3
95 - markdown=3.3.2=py38_0
96 - markupsafe=1.1.1=py38h8df0ef7_2
97 - matplotlib=3.3.4=py38h578d9bd_0
98 - matplotlib-base=3.3.4=py38h62a2d02_0
99 - mistune=0.8.4=py38h25fe258_1002

```

```

100 - mkl=2020.2=256
101 - mkl-service=2.3.0=py38he904b0f_0
102 - mkl_fft=1.3.0=py38h54f3939_0
103 - mkl_random=1.1.1=py38h0573a6f_0
104 - multidict=4.7.6=py38h7b6447c_1
105 - nbclassic=0.2.6=pyhd8ed1ab_0
106 - nbclient=0.5.3=pyhd8ed1ab_0
107 - nbconvert=6.0.7=py38h578d9bd_3
108 - nbformat=5.1.2=pyhd8ed1ab_1
109 - ncurses=6.2=he6710b0_1
110 - nest-asyncio=1.4.3=pyhd8ed1ab_0
111 - notebook=6.3.0=py38h578d9bd_0
112 - numpy=1.19.2=py38h54aff64_0
113 - numpy-base=1.19.2=py38hfa32c7d_0
114 - oauthlib=3.1.0=py_0
115 - olefile=0.46=pyh9f0ad1d_1
116 - openssl=1.1.1j=h27cfd23_0
117 - opt_einsum=3.1.0=py_0
118 - packaging=20.9=pyh44b312d_0
119 - pandas=1.1.3=py38he6710b0_0
120 - pandoc=2.12=h7f98852_0
121 - pandocfilters=1.4.2=py_1
122 - parso=0.8.1=pyhd8ed1ab_0
123 - pcre=8.44=he1b5a44_0
124 - pexpect=4.8.0=pyh9f0ad1d_2
125 - pickleshare=0.7.5=py_1003
126 - pillow=6.2.1=py38h6b7be26_0
127 - pip=20.2.4=py38_0
128 - prometheus_client=0.9.0=pyhd3deb0d_0
129 - prompt-toolkit=3.0.18=pyha770c72_0
130 - protobuf=3.14.0=py38h2531618_1
131 - pthread-stubs=0.4=h36c2ea0_1001
132 - ptyprocess=0.7.0=pyhd3deb0d_0
133 - pyasn1=0.4.8=py_0
134 - pyasn1-modules=0.2.8=py_0
135 - pycparser=2.20=py_2
136 - pygments=2.8.1=pyhd8ed1ab_0
137 - pyjwt=1.7.1=py38_0
138 - pyopenssl=19.1.0=py_1
139 - pyparsing=2.4.7=pyh9f0ad1d_0
140 - pyqt=5.9.2=py38h05f1152_4
141 - pyrsistent=0.17.3=py38h25fe258_1
142 - pysocks=1.7.1=py38_0
143 - python=3.8.5=h7579374_1
144 - python-dateutil=2.8.1=py_0
145 - python-flatbuffers=1.12=pyhd3eb1b0_0
146 - python_abi=3.8=1_cp38
147 - pytz=2021.1=pyhd8ed1ab_0
148 - pyyaml=5.3.1=py38h8df0ef7_1
149 - pyzmq=19.0.2=py38ha71036d_2
150 - qt=5.9.7=h5867ecd_1
151 - readline=8.0=h7b6447c_0
152 - requests=2.24.0=py_0
153 - requests-oauthlib=1.3.0=py_0
154 - rsa=4.6=py_0
155 - scikit-learn=0.23.2=py38h0573a6f_0
156 - scipy=1.6.1=py38h91f5cce_0
157 - send2trash=1.5.0=py_0

```

```

158 - setuptools=50.3.0=py38hb0f4dca_1
159 - sip=4.19.13=py38he6710b0_0
160 - six=1.15.0=py_0
161 - sniffio=1.2.0=py38h578d9bd_1
162 - sqlite=3.33.0=h62c20be_0
163 - tensorboard=2.4.0=pyhc547734_0
164 - tensorboard-plugin-wit=1.6.0=py_0
165 - tensorflow=2.4.1=gpu_py38h8a7d6ce_0
166 - tensorflow-base=2.4.1=gpu_py38h29c2da4_0
167 - tensorflow-estimator=2.4.1=pyheb71bc4_0
168 - tensorflow-gpu=2.4.1=h30adc30_0
169 - termcolor=1.1.0=py38_1
170 - terminado=0.9.3=py38h578d9bd_0
171 - testpath=0.4.4=py_0
172 - threadpoolctl=2.1.0=pyh5ca1d4c_0
173 - tk=8.6.10=hbc83047_0
174 - tornado=6.1=py38h25fe258_0
175 - traitlets=5.0.5=py_0
176 - urllib3=1.25.11=py_0
177 - wcwidth=0.2.5=pyh9f0ad1d_2
178 - webencodings=0.5.1=py_1
179 - werkzeug=1.0.1=py_0
180 - wheel=0.35.1=py_0
181 - wrapt=1.12.1=py38h7b6447c_1
182 - xorg-libxau=1.0.9=h14c3975_0
183 - xorg-libxdmcp=1.1.3=h516909a_0
184 - xz=5.2.5=h7b6447c_0
185 - yaml=0.2.5=h516909a_0
186 - yarl=1.6.2=py38h7b6447c_0
187 - zeromq=4.3.4=h2531618_0
188 - zipp=3.3.1=py_0
189 - zlib=1.2.11=h7b6447c_3
190 - zstd=1.4.5=h9ceee32_0
191 prefix: /home/gtu_hpc/anaconda3/envs/tensor

```



## 4 Algorithm Selection

The data contains one column for year(integer datatype) and 90 columns for timbre values(float datatype). It looks as follows:

	0	1	2	3	4	5	6	7	8	9	...	81	82	83	84
0	2001	49.94357	21.47114	73.07750	8.74861	-17.40628	-13.09905	-25.01202	-12.23257	7.83089	...	13.01620	-54.40548	58.99367	15.37344
1	2001	48.73215	18.42930	70.32679	12.94636	-10.32437	-24.83777	8.76630	-0.92019	18.76548	...	5.66812	-19.68073	33.04964	42.87836
2	2001	50.95714	31.85602	55.81851	13.41693	-6.57898	-18.54940	-3.27872	-2.35035	16.07017	...	3.03800	26.05866	-50.92779	10.93792
3	2001	48.24750	-1.89837	36.29772	2.58776	0.97170	-26.21683	5.05097	-10.34124	3.55005	...	34.57337	-171.70734	-16.96705	-46.67617
4	2001	50.97020	42.20998	67.09964	8.46791	-15.85279	-16.81409	-12.48207	-9.37636	12.63699	...	9.92661	-55.95724	64.92712	-17.72522
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
515340	2006	51.28467	45.88068	22.19582	-5.53319	-3.61835	-16.36914	2.12652	5.18160	-8.66890	...	4.81440	-3.75991	-30.92584	26.33968
515341	2006	49.87870	37.93125	18.65987	-3.63581	-27.75665	-18.52988	7.76108	3.56109	-2.50351	...	32.38589	-32.75535	-61.05473	56.65182
515342	2006	45.12852	12.65758	-38.72018	8.80882	-29.29985	-2.28706	-18.40424	-22.28726	-4.52429	...	-18.73598	-71.15954	-123.98443	121.26989
515343	2006	44.16614	32.38368	-3.34971	-2.49165	-19.59278	-18.67098	8.78428	4.02039	-12.01230	...	67.16763	282.77624	-4.63677	144.00125
515344	2005	51.85726	59.11655	26.39436	-5.46030	-20.69012	-19.95528	-6.72771	2.29590	10.31018	...	-11.50511	-69.18291	60.58456	28.64599

515345 rows x 91 columns

Figure 1: Dataset

Since the output class is a numeric discrete value, regression can be applied to this model.

For this dataset the following algorithms are used:

- Linear Regression using sklearn
- Single Layer NN using keras
- Multilayer Layer NN using keras

## 5 Neural Networks

A neural network is a series of algorithms that endeavors to recognize underlying relationships in a set of data through a process that mimics the way the human brain operates. It is capable of complex solutions which may not be possible with simpler algorithms. It best replicates the working of a human brain on a particular machine learning problem.

A single-layer neural network represents the most simple form of neural network, in which there is only one layer of input nodes that send weighted inputs to a subsequent layer of receiving nodes, or in some cases, one receiving node. This single-layer design was part of the foundation for systems which have now become much more complex.




Figure 2: Example of a Single Layer Neural Network

A Multi Layer Neural Network contains one or more hidden layers (apart from one input and one output layer). While a single layer perceptron can only learn linear functions, a multi layer perceptron can also learn non – linear functions.




Figure 3: Example of a Multi Layer Neural Network

### 5.1 Preprocessing

```

1 import pandas as pd
2 import numpy as np
3
4 df = pd.read_csv('YearPredictionMSD.txt', header=None)
5 df
6 for idx in range(X.shape[1]):
7     X[:,idx]=X[:,idx]/max(X[:,idx])
8
9 train_x=X[:463715]
10 test_x=X[463716:]
11 train_y=Y[:463715]
12 test_y=Y[463716:]
13
14 normalizer = preprocessing.Normalization()
15 normalizer.adapt(np.array(train_x))
16
17 first = np.array(train_x[:1])
18
19 with np.printoptions(precision=2, suppress=True):
20     print('First example:', first)
21     print()
22     print('Normalized:', normalizer(first).numpy())

```

```

First example: [[ 49.94  21.47  73.08   8.75 -17.41 -13.1  -25.01 -12.23   7.83
   -2.47   3.32  -2.32  10.21  611.11  951.09  698.11  408.98  383.71
  326.52  238.11  251.42  187.17  100.43  179.19   -8.42 -317.87   95.86
   48.1  -95.66 -18.06   1.97  34.42  11.73   1.37   7.79   -0.37
 -133.68 -83.26 -37.3   73.05 -37.37   -3.14 -24.22 -13.23  15.94
  -18.6   82.15  240.58 -10.29  31.58 -25.38   -3.91  13.29  41.55
   -7.26 -21.01  105.51  64.3   26.08 -44.59   -8.31   7.94 -10.74
  -95.45 -82.03 -35.59   4.7   70.96  28.09   6.02 -37.14 -41.12
   -8.41   7.2   -8.6   -5.91 -12.32  14.69 -54.32  40.15  13.02
  -54.41  58.99  15.37   1.11 -23.09  68.41  -1.82 -27.46   2.26]]

Normalized: [[ 1.08  0.39  1.83  0.47 -0.48 -0.28 -1.55 -1.31  0.39 -0.67  0.79 -0.58
 -1.06 -1.04 -0.8  -0.74 -1.05 -0.86 -0.87 -0.9  -0.67 -0.83 -1.01 -0.73
 -0.42 -0.5   0.26  0.35 -0.68 -0.47 -0.03  0.15  0.03  0.1  0.17 -0.68
 -0.2  -0.44  0.58  0.24 -0.3  -0.18  0.38 -0.43  0.42 -0.46  0.01  0.37
  0.36  0.05 -0.34 -0.43  0.01  0.48  0.05 -0.31  0.  0.24 -0.08 -0.11
 -0.19  0.15 -0.27  0.14 -0.37 -0.28  0.02  0.37 -0.04  0.19 -0.11 -0.2
  0.11  0.3   0.2  -0.01  0.04 -0.12  0.25  0.11 -0.08  0.11  0.14 -0.24
  0.05 -0.36  0.54 -0.47 -0.26  0.04]]

```

Figure 4: Normalized input

## 5.2 Model Preparation

Here we train the neural network and calculate the loss for it using the predict function using a single layer and multilayer neural network.

### 5.2.1 NN with keras

```

1 linear_model = tf.keras.Sequential([
2     normalizer,
3     layers.Dense(units=89)
4 ])
5
6 learning_rate=0.1
7 epoch=10

```

```

8 linear_model.compile(
9     optimizer=tf.optimizers.Adam(learning_rate=learning_rate),
10    loss='mean_absolute_error')
11 # mean_squared_error
12 history = linear_model.fit(
13     train_x, train_y,
14     epochs=epoch,
15     # suppress logging
16     verbose=0,
17     # Calculate validation results on 20% of the training data
18     validation_split = 0.2)
19 import matplotlib.pyplot as plt
20
21 def plot_loss(history):
22     plt.plot(history.history['loss'], label='loss')
23     plt.plot(history.history['val_loss'], label='val_loss')
24     plt.ylim([0, 10])
25     plt.xlabel('Epoch')
26     plt.ylabel('Error [MPG]')
27     plt.legend()
28     plt.grid(True)
29     plot_loss(history)
30
31 linear_model.evaluate(test_x, test_y, verbose=0)
32 test_predictions = linear_model.predict(test_x)
33 error = test_predictions - test_y
34 # Print MSE
35 sum(sum(np.square(error)))/2/len(test_y)

```

### 5.2.2 DNN with keras

```

1 learning_rate=0.1
2 epoch=10
3 def build_and_compile_model(norm):
4     model = keras.Sequential([
5         norm,
6         layers.Dense(90, activation='relu'),
7         # layers.Dense(180, activation='relu'),
8         layers.Dense(135, activation='relu'),
9         layers.Dense(91, activation='relu'),
10        layers.Dense(89)
11    ])
12
13    model.compile(loss='mean_absolute_error',
14        optimizer=tf.keras.optimizers.Adam(learning_rate))
15    return model
16 dnn_model = build_and_compile_model(normalizer)
17 dnn_model.summary()
18
19 history = dnn_model.fit(
20     train_x, train_y,
21     validation_split=0.2,
22     verbose=0, epochs=epoch)
23
24 import matplotlib.pyplot as plt
25
26 def plot_loss(history):

```

```


27     plt.plot(history.history['loss'], label='loss')
28     plt.plot(history.history['val_loss'], label='val_loss')
29     plt.ylim([0, 10])
30     plt.xlabel('Epoch')
31     plt.ylabel('Error [MPG]')
32     plt.legend()
33     plt.grid(True)
34     plot_loss(history)
35
36     dnn_model.evaluate(test_x, test_y, verbose=0)
37     test_predictions = dnn_model.predict(test_x)
38     error = test_predictions - test_y
39     # Print MSE
40     sum(sum(np.square(error)))/2/len(test_y)
41

```

## 5.3 Model Testing

Testing the neural network, we get the mean square error as follows:

```
plot_loss(history)
```



```
linear_model.evaluate(test_x, test_y, verbose=0)
```

```
0.6746880412101746
```

```
linear_model.metrics_names
```

```
['loss']
```

```
test_predictions = dnn_model.predict(test_x)
error = test_predictions - test_y
error
```

```

-----
NameError                                Traceback (most recent call last)
<ipython-input-18-5ccca7376bfd> in <module>
----> 1 test_predictions = dnn_model.predict(test_x)
      2 error = test_predictions - test_y
      3 error

NameError: name 'dnn_model' is not defined

```

```
sum(sum(np.square(error)))/2/len(test_y)
```

```
43.427506260108174
```

Figure 5: MSE for NN 0.1 learning rate and epoch=10

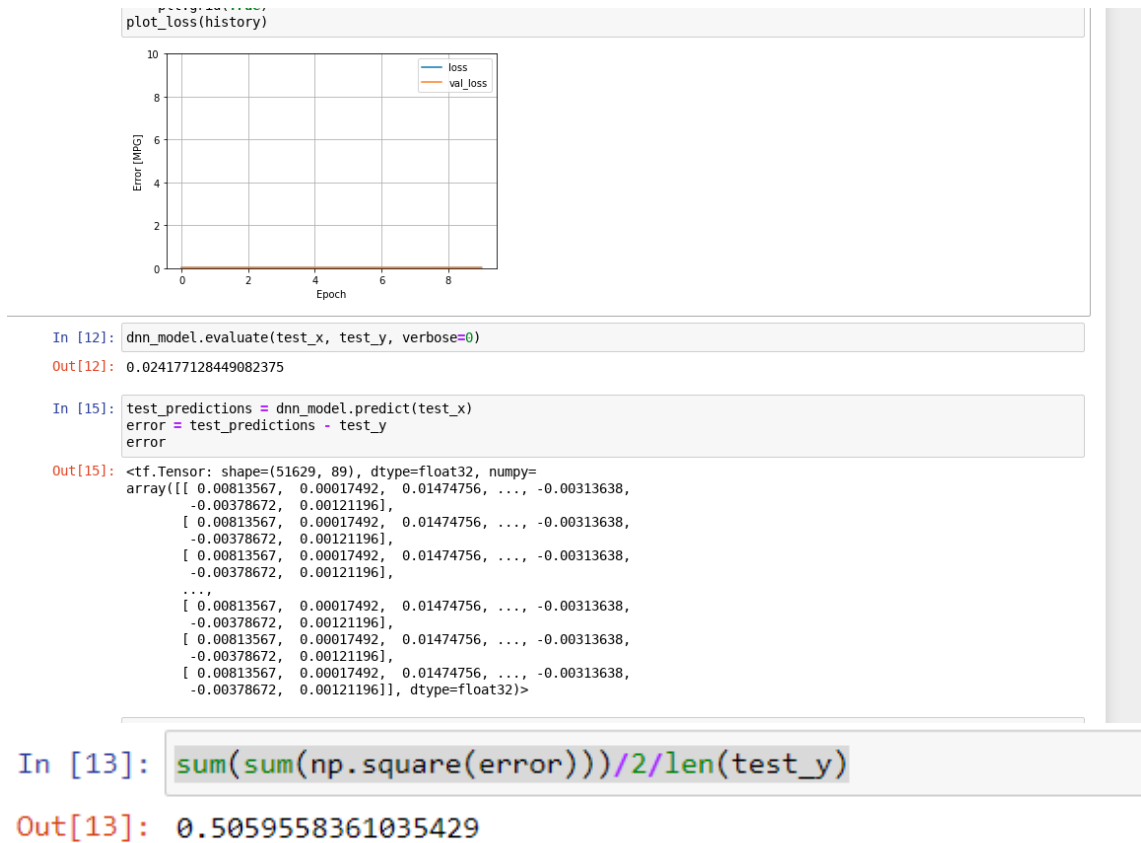


Figure 6: MSE for DNN 0.1 learning rate and epoch=10

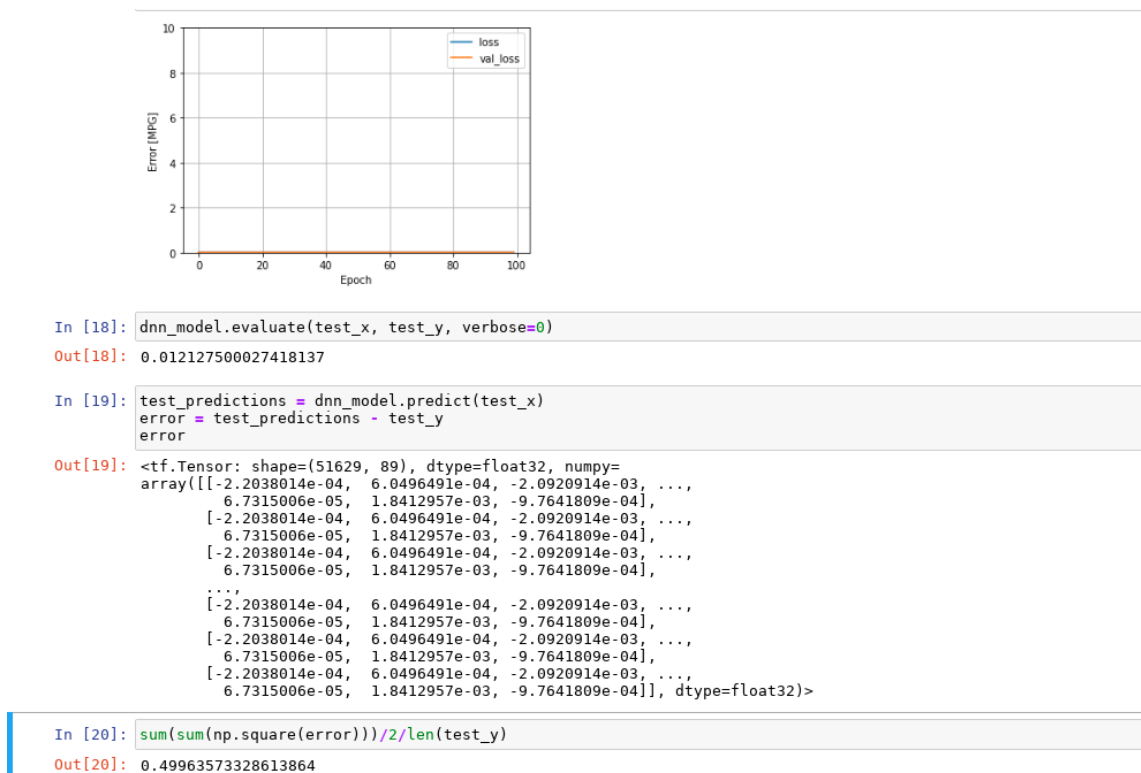


Figure 7: MSE for DNN 0.01 learning rate and epoch=100

## 6 Linear Regression

Linear regression attempts to model the relationship between two variables by fitting a linear equation to observed data. Hence, it tried to explain the relationship between the output Y and input(s) X by creating an equation for a line.




Figure 8: Example of a linear regression model fitting

### 6.1 Preprocessing

```
1 import pandas as pd
2 import numpy as np
3
4 df = pd.read_csv('YearPredictionMSD.txt', header=None)
5 df
6 for idx in range(X.shape[1]):
7     X[:,idx]=X[:,idx]/max(X[:,idx])
8
9 train_x=X[:463715]
10 test_x=X[463716:]
11 train_y=Y[:463715]
12 test_y=Y[463716:]
```

### 6.2 Model Preparation

```
1 from sklearn.linear_model import LinearRegression
2 model = LinearRegression().fit(train_x, train_y)
3
4 r_sq = model.score(train_x, train_y)
5
6 print('coefficient of determination:', r_sq)
7 print('intercept:', model.intercept_)
8 print('slope:', model.coef_)
```

### 6.3 Model Testing

```
1 y_pred = model.predict(test_x)
2 y_pred=np.around(y_pred)
3
```

```

4 print('predicted response:', y_pred, sep='\n')
5 print('actual response:', test_y, sep='\n')
6
7 from sklearn.metrics import mean_squared_error
8 mean_squared_error(test_y, y_pred)

```

```

print('coefficient of determination:', r_sq)
print('intercept:', model.intercept_)
print('slope:', model.coef_)

```

```

coefficient of determination: 0.3602828223068081
intercept: 1969.1274332562098
slope: [ 31.99683631 -9.77544382 -3.40378969 -15.39830478 16.25861626
-45.67552259 -3.19350038 -0.61672409  7.93201453 -0.35578293
-21.45137866 22.24266597 -32.78038255  5.23110072 -29.14189796
15.34648434 13.43595115  79.73408731 -21.87146977 102.39474909
-9.4270343 -51.82235668 -24.46285518 -1.06584318  8.74403999
-37.08875404  8.32614338 -16.06079312  58.73159  37.04968064
-25.80369195 -21.28523073 18.16078946  8.13672016 28.19854803
 0.7209632 -13.47509332  3.49620344 -24.71320915 37.21186424
23.03969494 39.21651129  2.51669537 19.29750877 63.92556694
-5.88159329  8.92052974 -12.21537432 26.34829118 -25.00555365
-8.79449499  2.49277095 -61.0518601 75.71623298 -25.60892549
-30.59776697 19.46240851 -0.53951436 -12.79946087  9.97701374
 0.17772797 -6.683211  4.82130752  9.7429527 -35.36484008
-19.1143826 -18.48040118 32.95391027 -7.88734508 16.5522544
-5.7644389 -6.25066557 -16.73583417 -36.82316118 30.25973589
10.3951017 -1.90816456 -36.19017985 26.99644702 -0.75357372
-0.51186282  0.87699778 -8.00938572 -8.36472649 -13.21635156
 0.41984244]

```

```

from sklearn.metrics import confusion_matrix, mean_squared_error
mean_squared_error(test_y, y_pred)

```

738.25

```
len([0 for ele in test_y-y_pred if ele==0])/len(y_pred)*100
```

5.0

Figure 9: Linear Regression params and error



## 7 Conclusion

Comparing the mean-sqaure error term of the above alogorithms we get:

Algorithm	MSE
NN with learning rate=0.1 and epoch=10	43.427
DNN with learning rate=0.1 and epoch=10	0.5059
DNN with learning rate=0.1 and epoch=100	0.4996
Linear Regression	738.25

We see that DNN is by far a better option when selecting the algorithm, as it has the least mean square error amongst the other models.

Multilayer networks solve the classification problem for non linear sets by employing hidden layers, whose neurons are not directly connected to the output. The additional hidden layers can be interpreted geometrically as additional hyper-planes, which enhance the separation capacity of the network. Hence, we can conclude that DNN(multi layered neural networks) is more suitable for regression problems whose solution is not linear and with many features.